



A Internal validation

5-fold cross-validation · N = 1238



B Temporal validation

fitted on 2016-2023 · N = 413



One-versus-rest ROC curves with 95% bootstrap confidence intervals (2000 resamples of cases); class sizes in parentheses. Panel A imputes separately within each fold; panel B applies the 2016-2023 model unchanged. Apparent (in-sample) curves are deliberately not shown — they overstate discrimination by about 0.03 macro AUC. Surgery separates sharply in both panels; the observation-endoscopy boundary does not, which we argue is a property of the decision rather than of the model.