




One imputation draw per point (not pooled multiple imputation); shaded bands are ± 1 SD across 10 repeats per training size. Surgery separates with little training data; observation and endoscopy are the comparison this curve is meant to inform — see `outputs/learning_curve.txt` for the head-vs-tail slope test and its bearing on the manuscript's central claim.